ORACLE

Deploy Prometheus and Grafana on OCI Container Engine for Kubernetes

Feb 2023, Version 1.1

Desmond Muriu
EMEA Cloud Specialist Engineer - Compute

Table of contents

NB: **Avoid directly doing copy/paste from this document since it could include hidden characters resulting into command lines failures.**

1. Log In to OCI

Check this blog post on the details of How to access your OCI console -> https://docs.oracle.com/en-us/iaas/Content/GSG/Tasks/signingin.htm

2. Create Cluster

1. From the OCI Services menu (top left hamburger button), click Developer Services > Kubernetes Clusters (OKE).

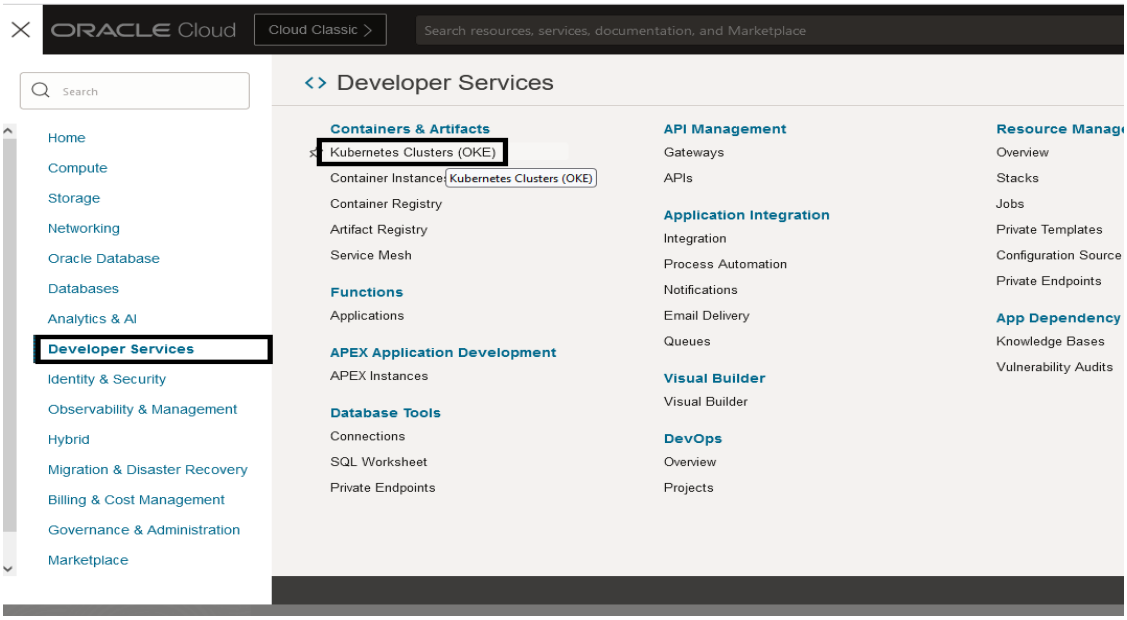

2. Under List Scope, select the compartment in which you would like to create a cluster then click create Cluster

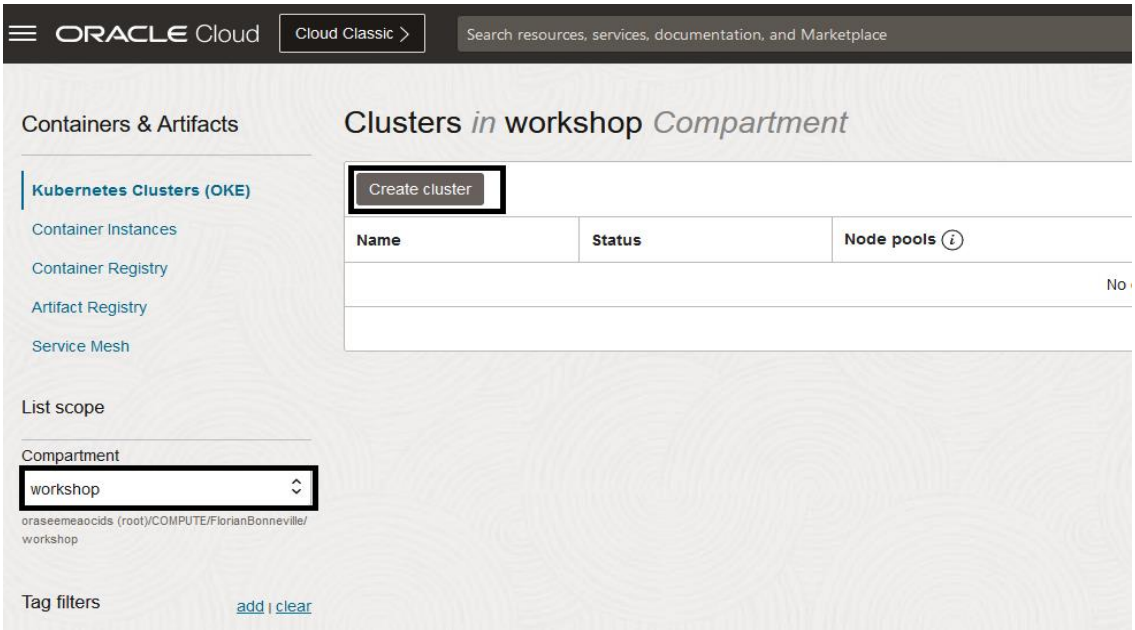

3. On the next pop up, choose Quick Create and click Submit.

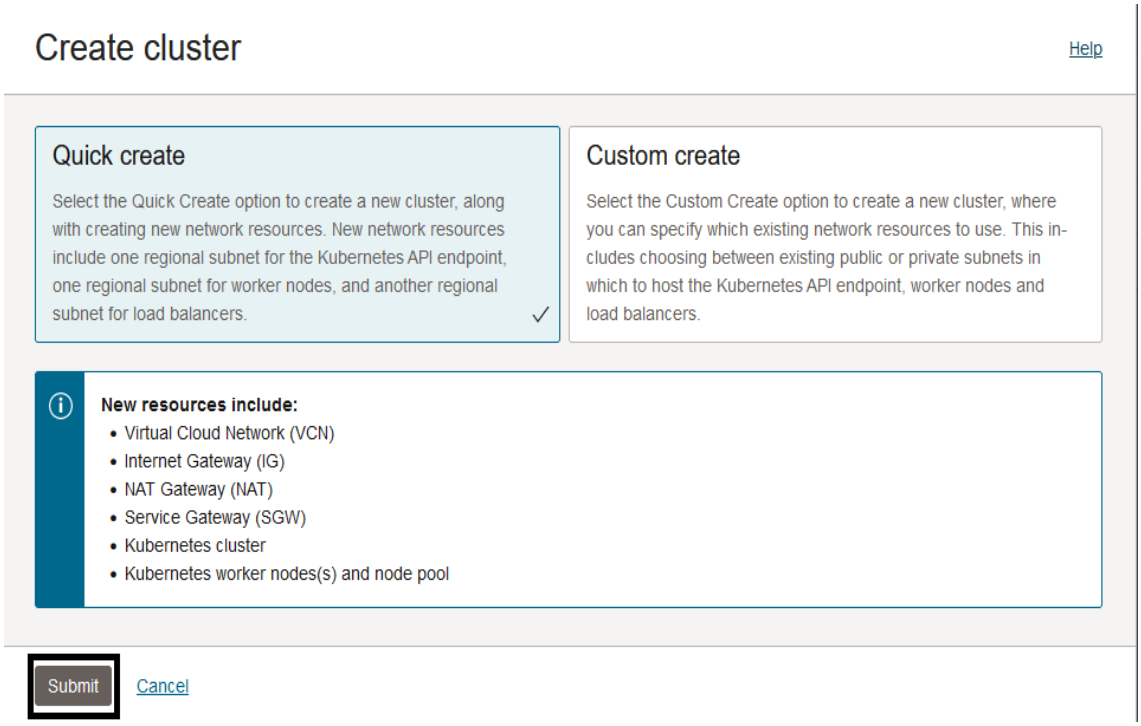

4. Fill out the Cluster Details on the Form that pops up .i.e.

- **Name**: Provide a name
- **Compartment**: Choose your compartment
- **Kubernetes Version**: Choose the most recent version(At the time of making this tutorial, the most recent version is v1.25.4)
- **Kubernetes API Endpoint**: Public Endpoint(this will allow cluster access on OCI Cloud shell)
- **Node Type**: Managed
- **Kubernetes Worker Nodes**: Private Workers
- **Shape and Image**: Select a pod shape, Number of OCPUS, Amount of RAM and Image based on your requirement (VM.Standard.E4.Flex, 2 OCPU and 8GB Ram will be used in this example)
- **Node count**: Provide number of nodes(3 in this example)

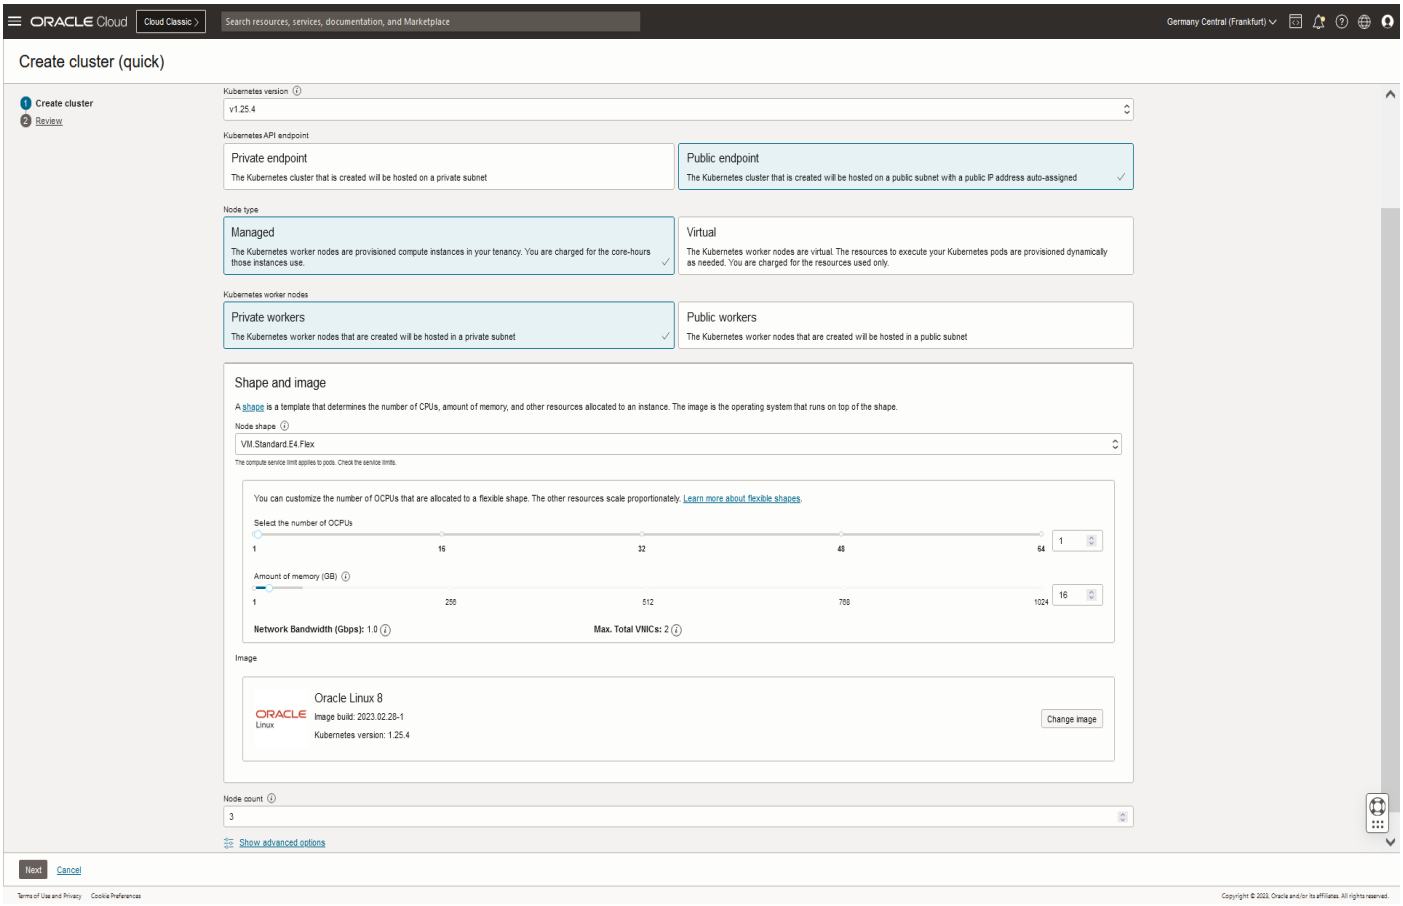

5. Click Next. Review the cluster details and finally create cluster.

3. Accessing the Cluster

You can use the Kubernetes command line tool **kubectl** to perform operations on a cluster you've created with Container Engine for Kubernetes. You can use the kubectl installation included in **OCI Cloud Shell**, or you can use a local installation of kubectl. In both cases, before you can use kubectl to access a cluster, you have to specify the cluster on which to perform operations by setting up the cluster's kubeconfig file.

In this tutorial OCI cloud shell will be used to interact with the OKE cluster.

To access the cluster from the cloud shell follow the below.

1. Select the deployed cluster from the Clusters list page.

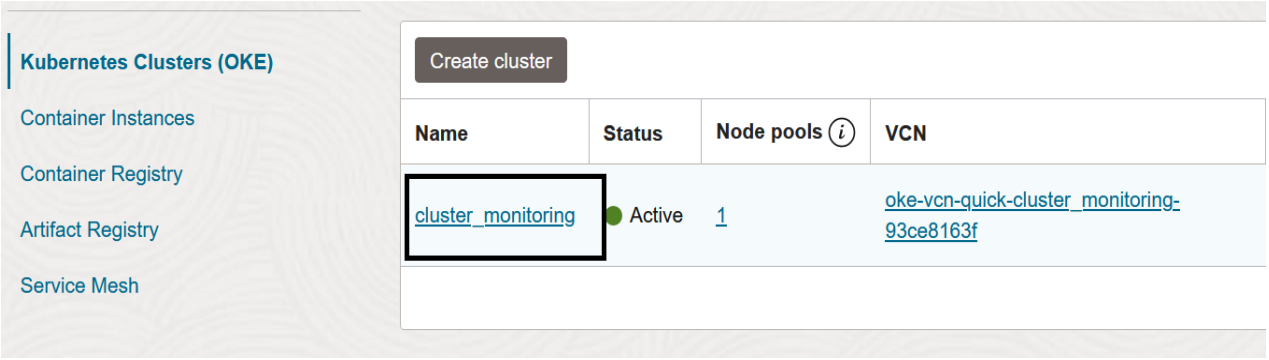

2. Click on the **Access Cluster** Link

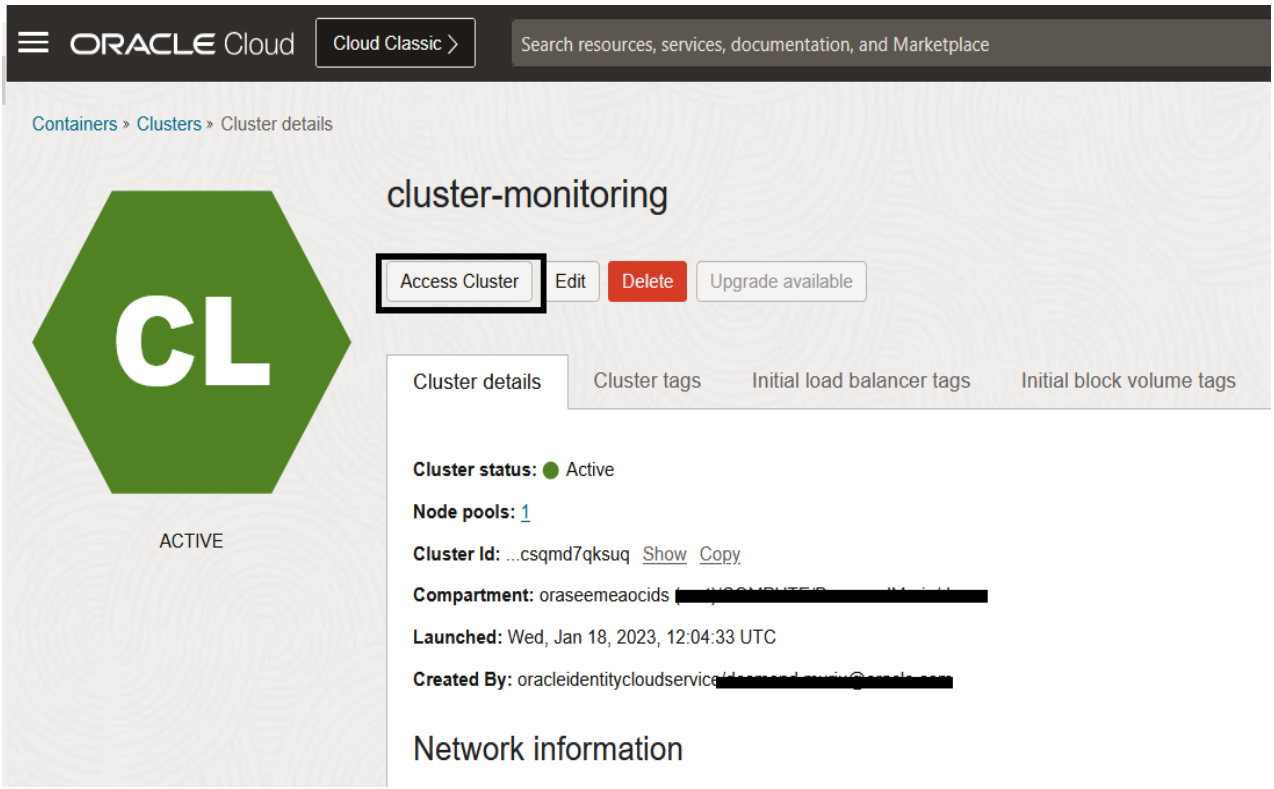

3. Launch **Cloud shell** from the top right of the OCI console and copy the command to access the cluster.

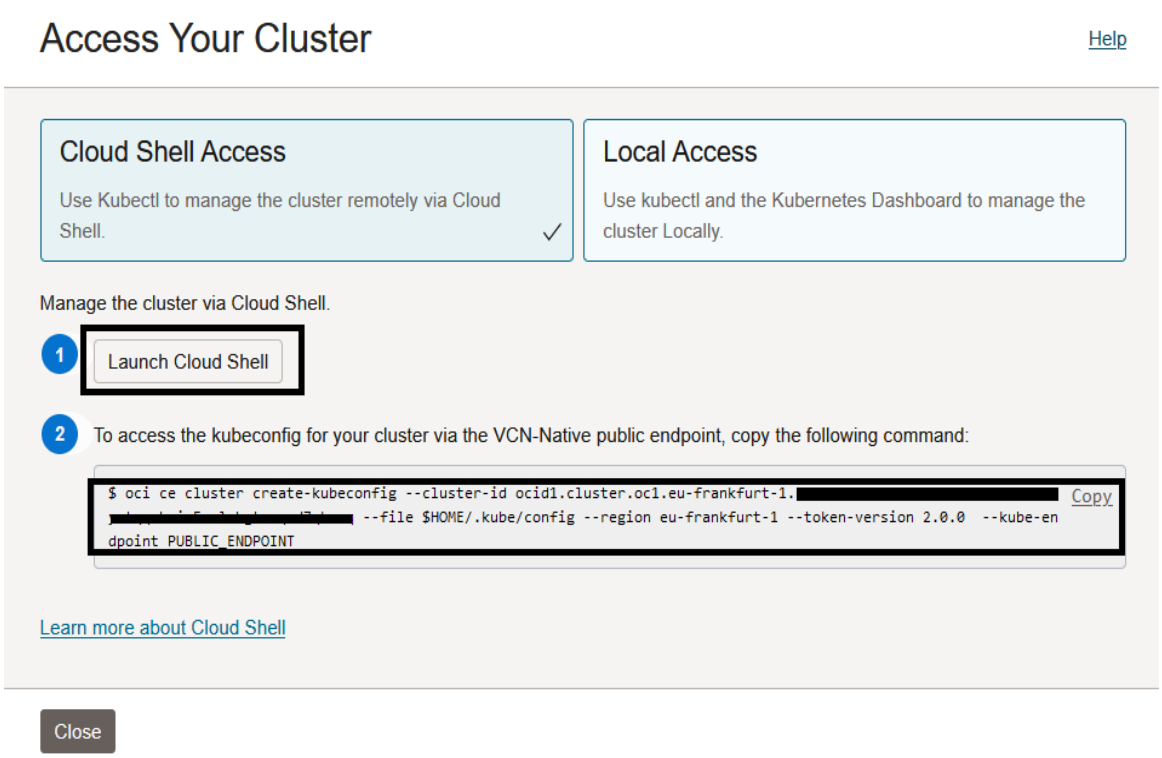

4. Enter a simple kubectl command to check you have access to the cluster eg **kubectl get all**

4. Install Prometheus and Grafana

1. Create monitoring namespace. A namespace could be thought of as a virtual cluster. It provides a means of organizing a group of resources within single cluster. In this case, we create a monitoring namespaces where all monitoring resources can be isolated in.

   kubectl create namespace monitoring

```
[redacted]@cloudshell:~ (eu-frankfurt-1)$ kubectl create namespace monitoring
namespace/monitoring created
```

2. Add helm repos

   helm repo add prometheus-community https://prometheus-community.github.io/helm-charts

   helm repo update

```
[redacted]@cloudshell:~ (eu-frankfurt-1)$ helm repo add prometheus-community https://prometheus-community.github.io/helm-char
"prometheus-community" already exists with the same configuration, skipping
[redacted]@cloudshell:~ (eu-frankfurt-1)$ helm repo update
Hang tight while we grab the latest from your chart repositories...
...Successfully got an update from the "ingress-nginx" chart repository
...Successfully got an update from the "jetstack" chart repository
...Successfully got an update from the "rancher-stable" chart repository
...Successfully got an update from the "prometheus-community" chart repository
Update Complete. ⎈Happy Helming!⎈
```

3. Install chart in the monitoring namespace

   helm install oke-prom --namespace monitoring prometheus-community/kube-prometheus-stack

```
[redacted]@cloudshell:~ (eu-frankfurt-1)$ helm install oke-prom --namespace monitoring  prometheus-community/kube-prometheus-

NAME: oke-prom
LAST DEPLOYED: Wed Feb  8 15:14:57 2023
NAMESPACE: monitoring
STATUS: deployed
REVISION: 1
NOTES:
kube-prometheus-stack has been installed. Check its status by running:
  kubectl --namespace monitoring get pods -l "release=oke-prom"

Visit https://github.com/prometheus-operator/kube-prometheus for instructions on how to create & configure Alertmanager and Promet
heus instances using the Operator.
```

4. Expose the grafana service to allow for external access.

   To get the service name use the *kubectl get svc –n monitoring* to check the list of services. The current service is a clusterIP type meaning it cannot be accessed outside of the cluster. The next task will be to edit this service to allow for external access via an OCI native Load balancer.

```
[redacted]@cloudshell:~ (eu-frankfurt-1)$ kubectl get svc -n monitoring
NAME                                      TYPE        CLUSTER-IP      EXTERNAL-IP   PORT(S)                      AGE
alertmanager-operated                     ClusterIP   None            <none>        9093/TCP,9094/TCP,9094/UDP   5m16s
oke-prom-grafana                          ClusterIP   10.96.90.24     <none>        80/TCP                       5m24s
oke-prom-kube-prometheus-s-alertmanager   ClusterIP   10.96.239.124   <none>        9093/TCP                     5m24s
oke-prom-kube-prometheus-s-operator       ClusterIP   10.96.12.254    <none>        443/TCP                      5m24s
oke-prom-kube-prometheus-s-prometheus     ClusterIP   10.96.94.93     <none>        9090/TCP                     5m24s
oke-prom-kube-state-metrics               ClusterIP   10.96.1.10      <none>        8080/TCP                     5m24s
oke-prom-prometheus-node-exporter         ClusterIP   10.96.250.197   <none>        9100/TCP                     5m24s
prometheus-operated                       ClusterIP   None            <none>        9090/TCP                     5m16s
```

   To do this run the following command:

   kubectl edit svc oke-prom-grafana -n monitoring

   Under the **annotations** section add *oci.oraclecloud.com/load-balancer-type: "lb"*

```
# Please edit the object below. Lines beginning with a '#' will be ignored,
# and an empty file will abort the edit. If an error occurs while saving this file will be
# reopened with the relevant failures.
#
apiVersion: v1
kind: Service
metadata:
  annotations:
    meta.helm.sh/release-name: oke-prom
    meta.helm.sh/release-namespace: monitoring
    oci.oraclecloud.com/load-balancer-type: "lb"
  creationTimestamp: "2023-02-08T15:15:17Z"
  labels:
    app.kubernetes.io/instance: oke-prom
    app.kubernetes.io/managed-by: Helm
    app.kubernetes.io/name: grafana
    app.kubernetes.io/version: 9.3.6
    helm.sh/chart: grafana-6.50.7
  name: oke-prom-grafana
  namespace: monitoring
  resourceVersion: "6473"
  uid: 15d863b5-1e9e-4f57-8213-5925fccc67e4
spec:
  clusterIP: 10.96.90.24
  clusterIPs:
  - 10.96.90.24
```

Also change the **type** from ClusterIp to Loadbalancer at the bottom of the config file

```
    app.kubernetes.io/version: 9.3.6
    helm.sh/chart: grafana-6.50.7
  name: oke-prom-grafana
  namespace: monitoring
  resourceVersion: "6473"
  uid: 15d863b5-1e9e-4f57-8213-5925fccc67e4
spec:
  clusterIP: 10.96.90.24
  clusterIPs:
  - 10.96.90.24
  internalTrafficPolicy: Cluster
  ipFamilies:
  - IPv4
  ipFamilyPolicy: SingleStack
  ports:
  - name: http-web
    port: 80
    protocol: TCP
    targetPort: 3000
  selector:
    app.kubernetes.io/instance: oke-prom
    app.kubernetes.io/name: grafana
  sessionAffinity: None
  type: Loadbalancer
status:
  loadBalancer: {}
-- INSERT --
```

After the update view the services again and this time note that the type will change to Loadbalancer and an external IP will also be provided. The Loadbalancer createad is a native OCI Load balancer which can also be viewed from the OCI console by going to OCI Services menu (top left hamburger button) > Networks > Loadbalancers

```
[redacted]@cloudshell:~ (eu-frankfurt-1)$ kubectl get svc -n monitoring
NAME                                      TYPE           CLUSTER-IP      EXTERNAL-IP      PORT(S)                      AGE
alertmanager-operated                     ClusterIP      None            <none>           9093/TCP,9094/TCP,9094/UDP   20m
oke-prom-grafana                          LoadBalancer   10.96.90.24     141.144.252.218  80:32001/TCP                 20m
oke-prom-kube-prometheus-s-alertmanager   ClusterIP      10.96.239.124   <none>           9093/TCP                     20m
oke-prom-kube-prometheus-s-operator       ClusterIP      10.96.12.254    <none>           443/TCP                      20m
oke-prom-kube-prometheus-s-prometheus     ClusterIP      10.96.94.93     <none>           9090/TCP                     20m
oke-prom-kube-state-metrics               ClusterIP      10.96.1.10      <none>           8080/TCP                     20m
oke-prom-prometheus-node-exporter         ClusterIP      10.96.250.197   <none>           9100/TCP                     20m
prometheus-operated                       ClusterIP      None            <none>           9090/TCP                     20m
```

The external-ip will be the endpoint to access the Grafana UI

5. Access Grafana UI

1. Enter the external IP on your browser to get access to the grafana UI.

   Default username: **admin**

   Default Password: **prom-operator(remember to change password after login)**

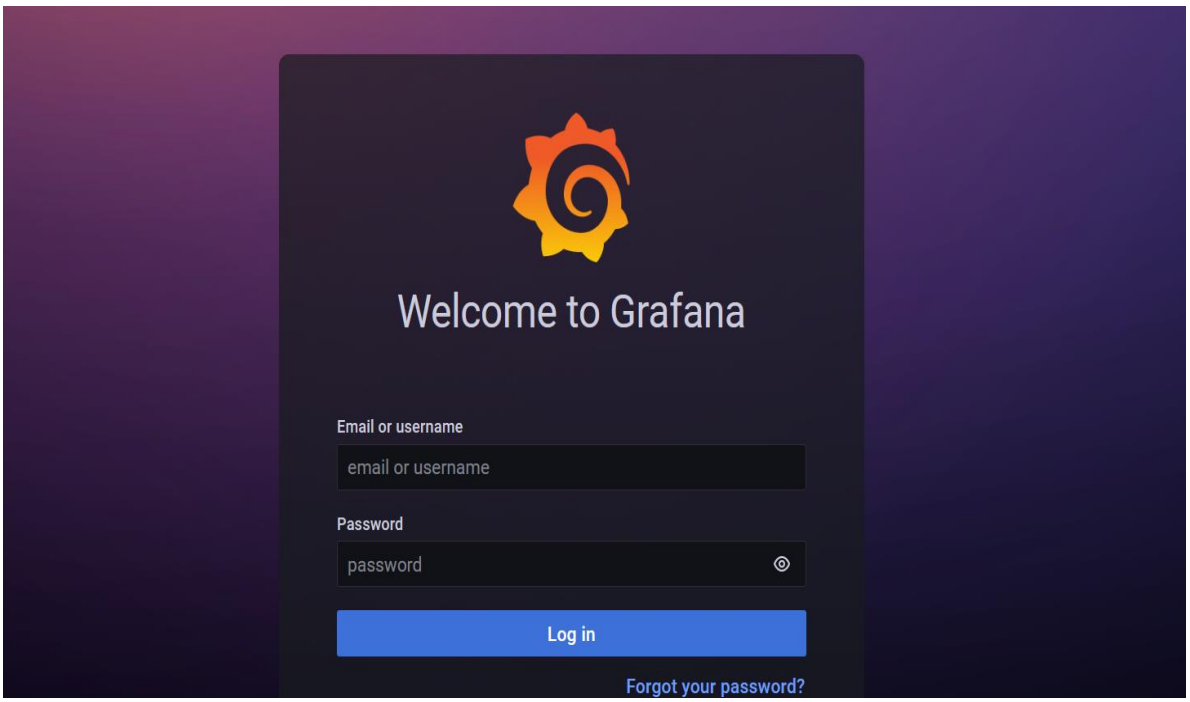

2. Changing Password. This can be done by clicking the bottom left Admin icon and selecting the **change password** option.

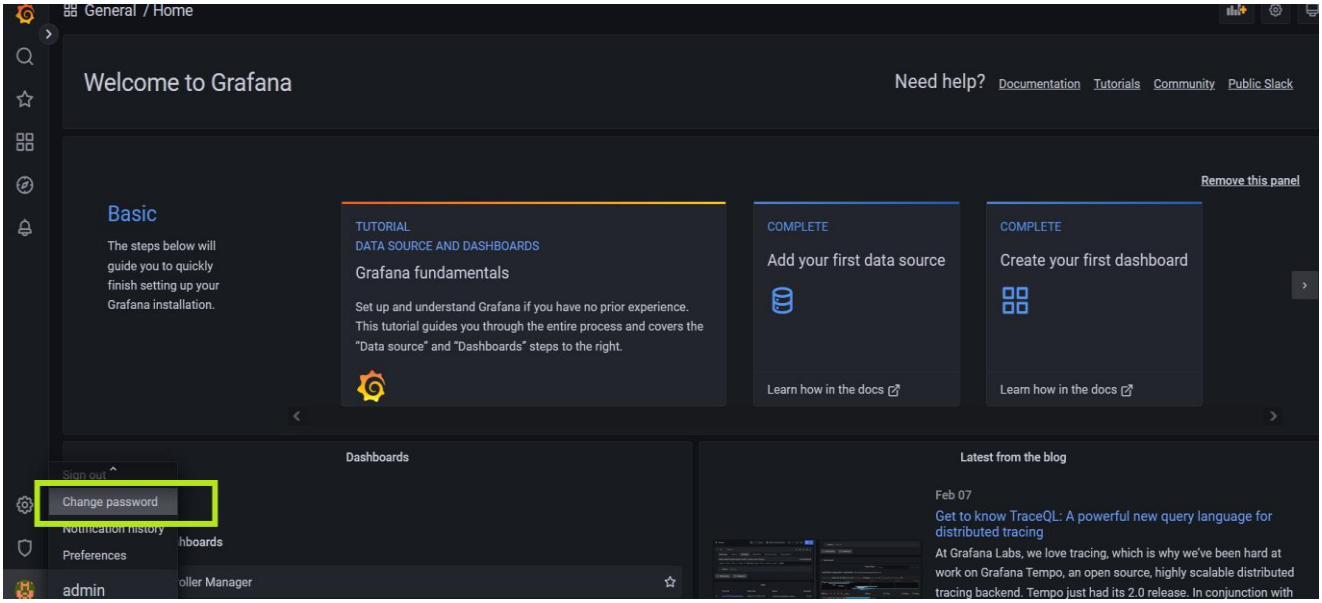

3. View **Data Sources** by selecting it from the **Configuration** options on the bottom left bar.

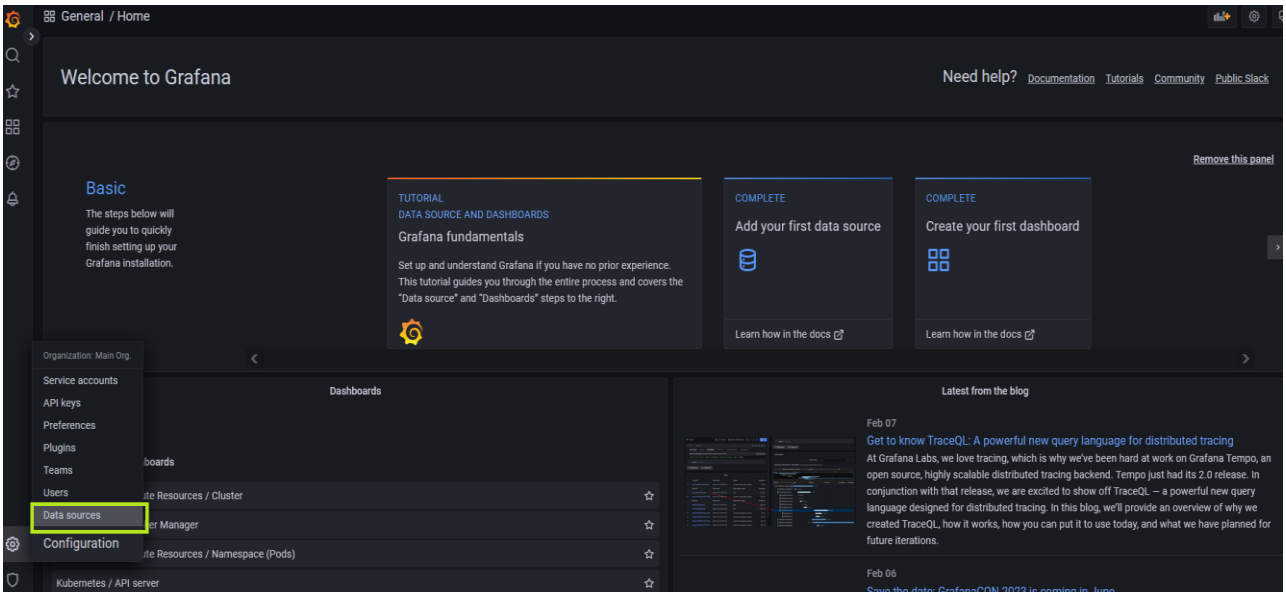

By default Prometheus is already configured.

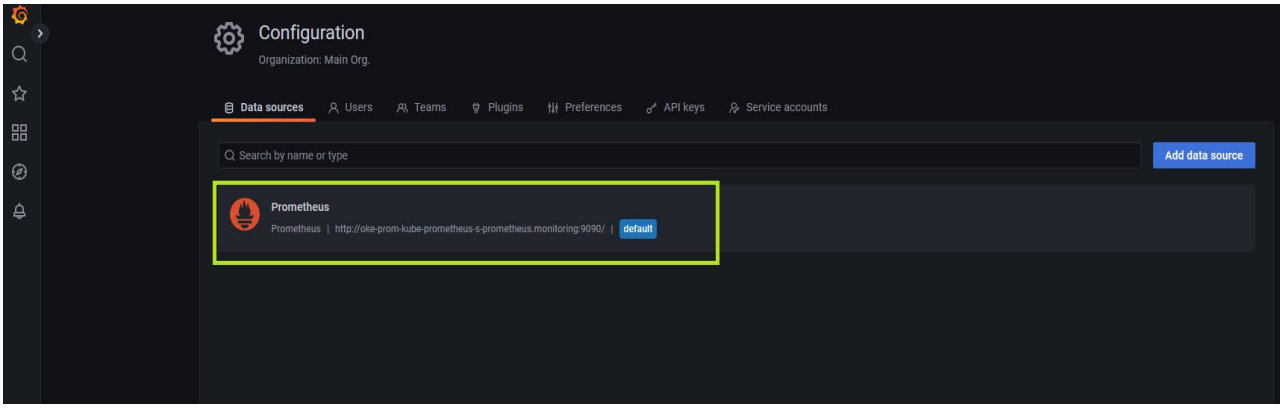

4. Accessing the default dashboards. This is achieved by clicking **Browse** from the **Dashboards** tab on the left

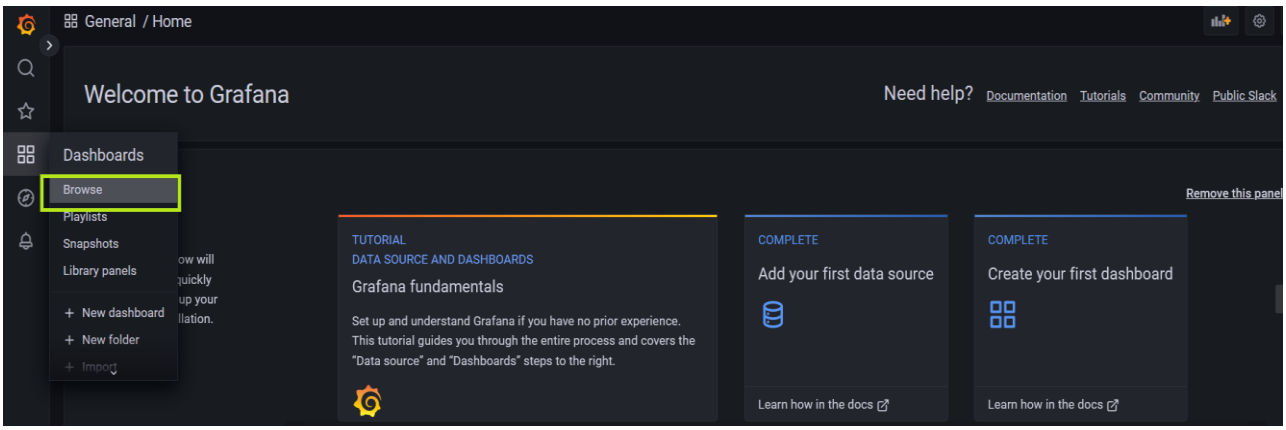

Eg. To view a summary of compute resources on your cluster , select **Kubernetes/Compute Reources/Cluster** option